


.NET - Programowanie równoległe ~~i trochę rozproszone~~ kiedy (i jak)

Tomasz Kopacz

35 minut

Demos: <https://github.com/tkopacz/msts2024-net-parallel>



"Standard" computers... (and servers)

2005

Pentium 4, HT, 3.74 GHz
512 MB, 1GB SOMETIMES
10GFLOPS | 10K MIPS

Xeon, 2 cores+ HT, 3.6 GHz
4 GB
1 048 576 4KB records

2015

i5, 6 cores + HT, 3.2 GHz
8 GB
1TFLOPS | 100K MIPS

E5, 8 cores + HT, 2.4 GHz
32 GB
8 388 608 4KB records

2024

I7/i9, 20 + HT, 3.4-5 GHz
32/128 GB
100TFLOPS | 1M MIPS

Gold, 28 cores + HT, 2.6 GHz
256 GB
67 108 864 4KB records
IN RAM!
67M (256 is not too much!)

*3.5 GHz = 1 cycle every 1/3 500 000 000 second = 2.85 ns
Speed of light = 3 * 10⁸ m / sec. 30 cm in 1 nanosecond!*

*Imagine – 100GHz, cycle 0.01ns, size 0.3 mm
("slowly" quantum physics)*

Business – real one

News webpage/sec

Wikipedia: 30K / sec, 300 machines, 100/machine

Twitter: 800 connection per second

NYT: 200-300 per sec, 10K peak

Orders/day

Restaurant: 50 (EU) - 2000 (China) per day
Grocery, Retail: 1000

E-commerce: 5000

Estimations!

Amazon: 33M

200M Prime Members

Allegro: 0.7M

Search, scale events

There are 0.6 million Google searches per second

100K/sec in Azure Event Hub (IoT, telemetry etc)

But:

Intelligent Thermometer
1 reading per 30 seconds

Hardware....

PCI-E

3.0: 15.8 GB/sec

4.0: 31.5 GB/sec

5.0: 63 GB/sec

(7.0,224) 242 GB/s

NVME: 3-12GB / sec

SSD/SATA: 0.5GB / sec

DDR5: 64GB/s

ETH 10Gbit: 1 GB / sec

8,112,077,412 people on Earth
 $\approx 6 \cdot 10^9$ mobiles + $3 \cdot 10^9$ „PC“
+ traffic from bots 😊



Clemens Vasters • 1st
Principal Architect, Messaging Services and Standards, Microsoft

2d (edited) ***

Azure Event Hubs handles roughly 120 million requests per second. No typo. We have quite a few customers who do more than 100K/sec just at the messaging level.

Edit: Made a correction to 1/7th of what I had originally written since I looked up into wrong column -- weekly instead of daily -- and we usually don't calculate req/sec at that level. Sorry. The number is still huge :)

Options to run paralel code in .NET

(normal Threads)

(async/await)

Tasks and Task Parallel Library

Parallel.For | Parallel.ForEach | await foreach (IAsyncEnumerable)

PLINQ (Data Parallelism)

Data Flow

Reactive Extension (Rx)

Scenarios

Server-side, console

Also:

One Blazor

One Windows Forms

What are the limits for concurrent programming

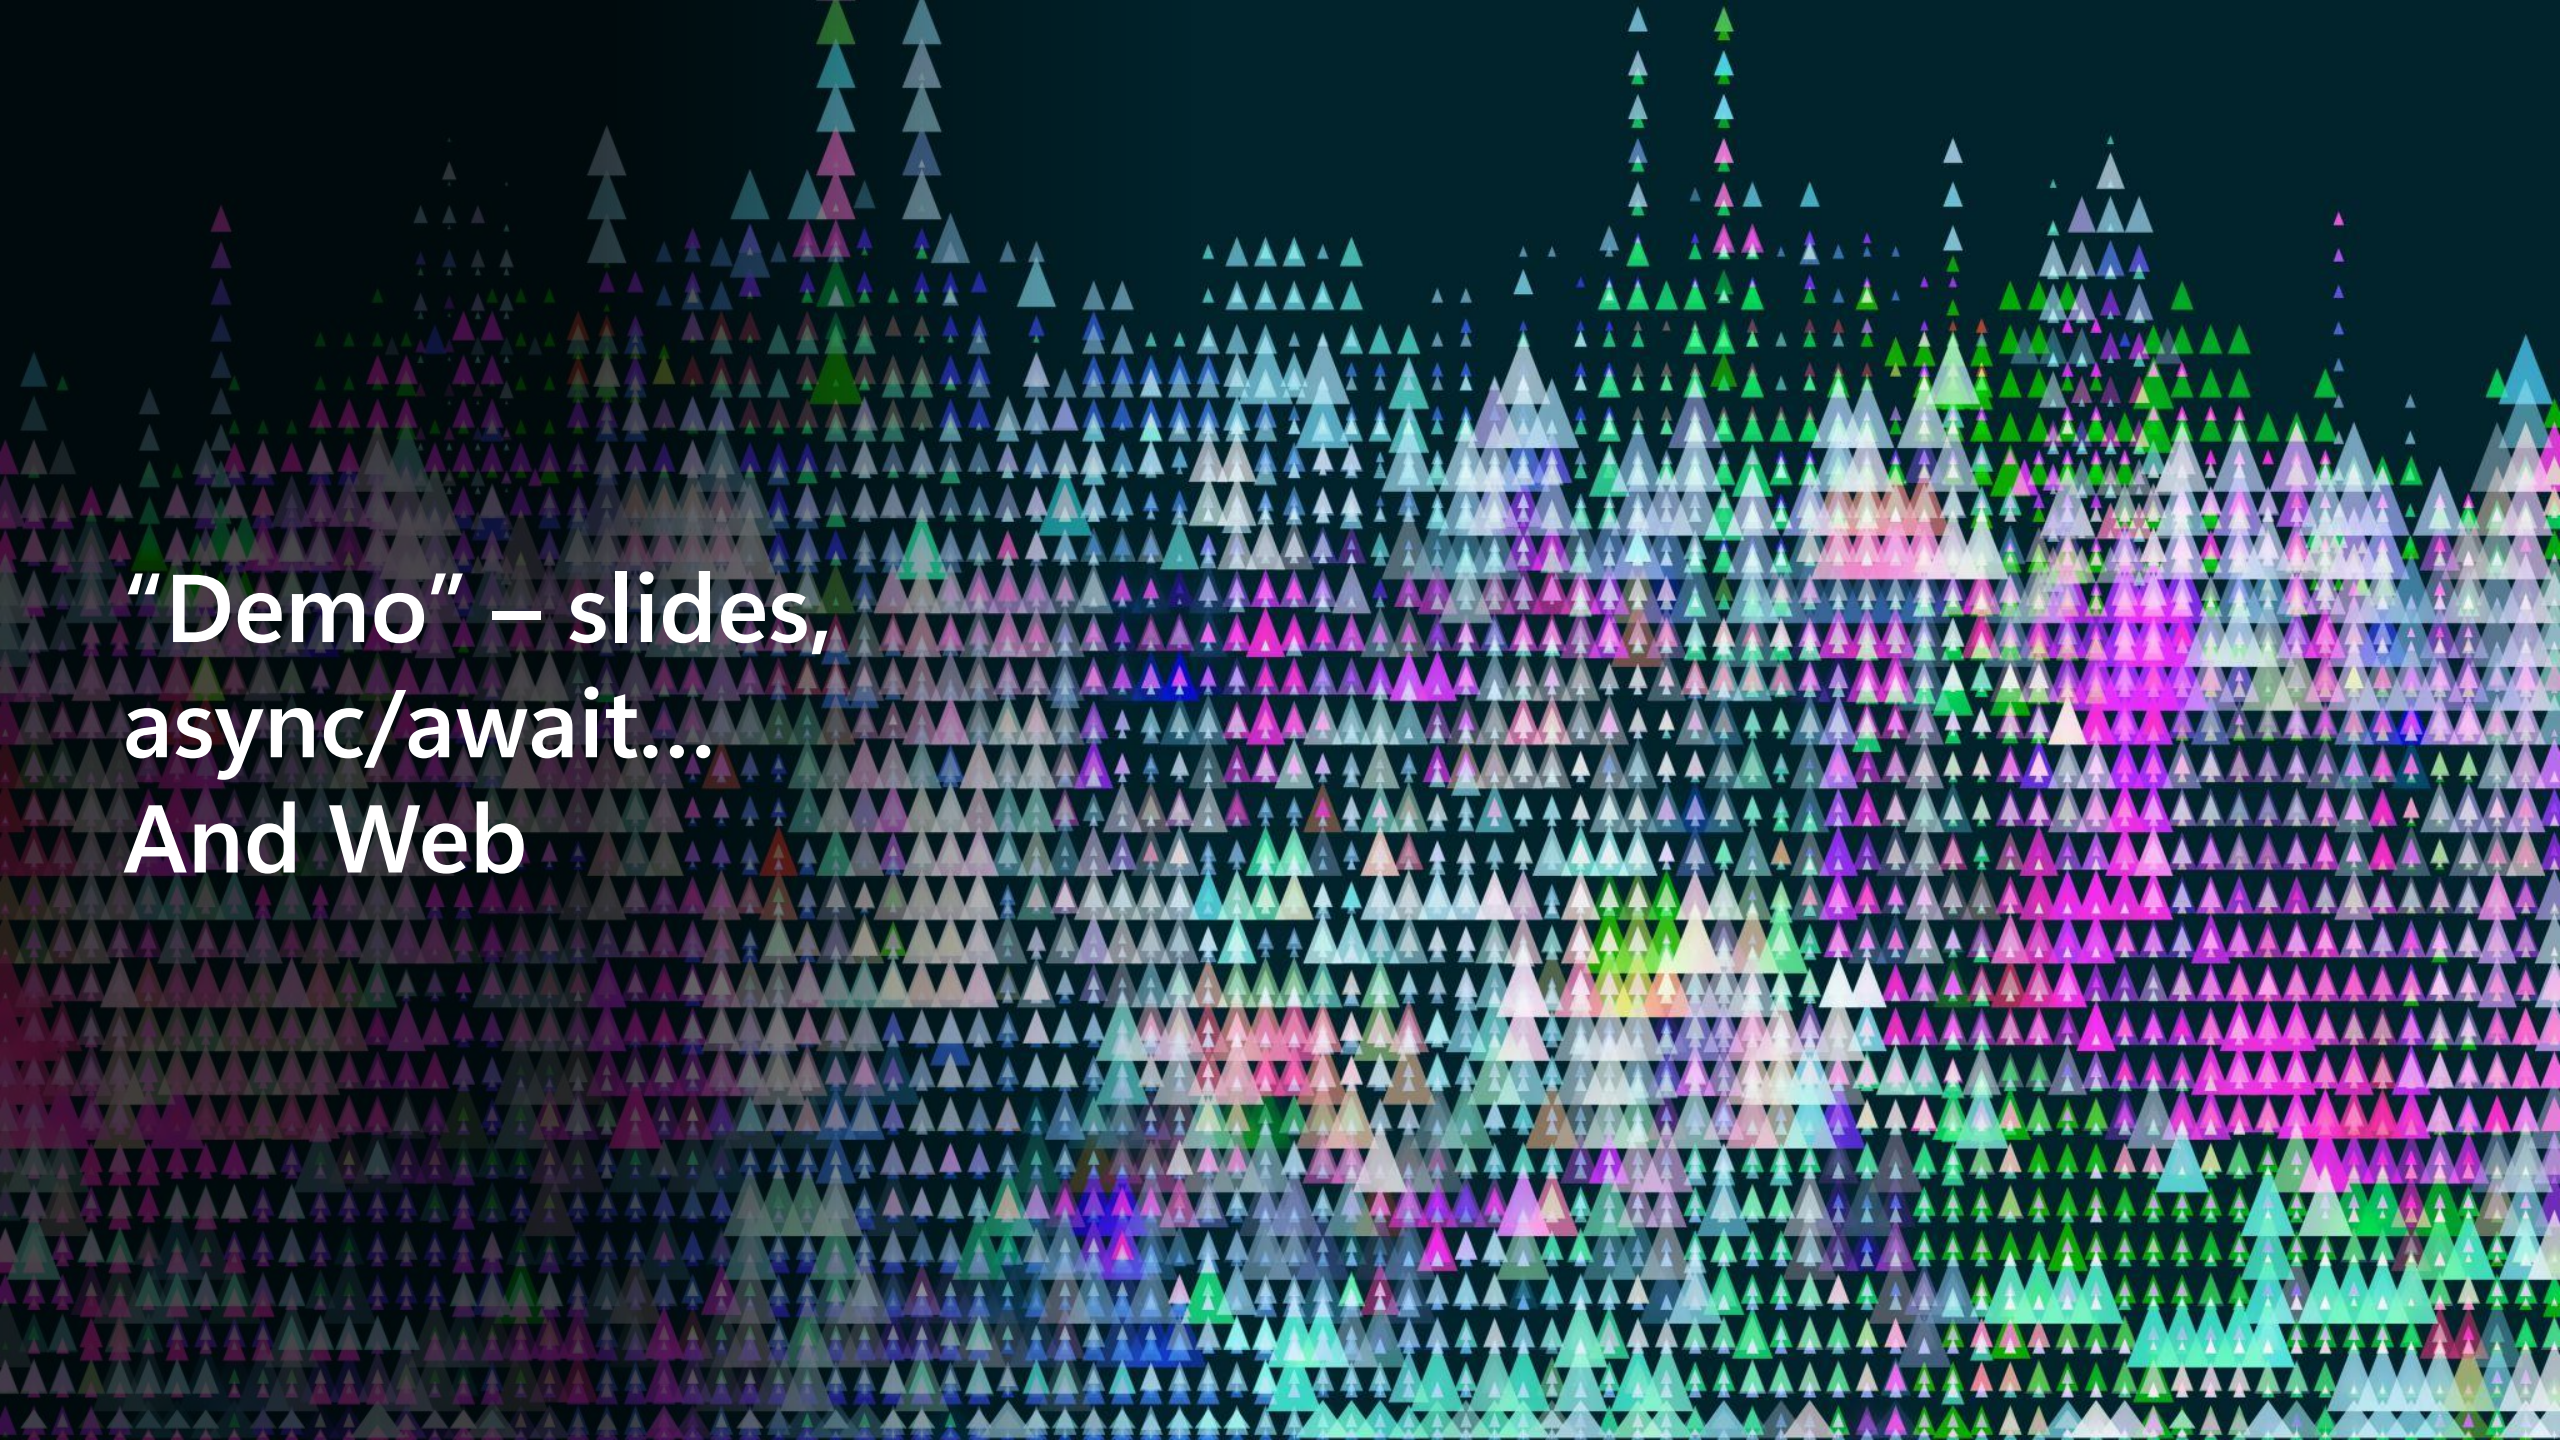
Local Hardware

- CPU/Cores/Threads
- Sockets number
- I/O throughput
- Memory (rare)

Remote (external) services

- SQL DTU, Max RU/s (CosmosDB), Tokens/second....
- Network throughput

Our own mind / knowledge / skills!



“Demo” – slides,
async/await...
And Web


```

app.MapGet("/long", () =>
{
    var forecast = Enumerable.Range(1, 500).Select(index =>
        new LongRecord
        (
            new string('t', index * 2 + Random.Shared.Next(255)),
            new string('d', index * 3 + Random.Shared.Next(255)),
            1000 + Random.Shared.Next(2000)
        ))
        .ToArray();
    return forecast;
});

```

Sync Web Api

Async Web Api

```

app.MapGet("/longasync", async () =>
{
    return await Task.Run(() =>
    {
        var forecast = Enumerable.Range(1, 500).Select(index =>
            new LongRecord
            (
                new string('t', index * 2 + Random.Shared.Next(255)),
                new string('d', index * 3 + Random.Shared.Next(255)),
                1000 + Random.Shared.Next(2000)
            ))
            .ToArray();
        return forecast;
    });
});

```

Numbers... (ACA, 1vCPU + 1GB, all defaults)

400 vUsers

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	4000	0	0.00%	3666.02	103	29388	2712.00	7291.80	9576.25	13189.95	79.62	60541.72	16.95
Sync	4000	0	0.00%	3666.02	103	29388	2712.00	7291.80	9576.25	13189.95	79.62	60541.72	16.95

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	4000	0	0.00%	3995.92	369	27644	3294.00	7281.60	8832.75	12063.70	70.57	53660.71	14.06
Async	4000	0	0.00%	3995.92	369	27644	3294.00	7281.60	8832.75	12063.70	70.57	53660.71	14.06

450 vUsers, Async – **2x** more errors than sync (and will grow **FASTER**)

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	4500	19	0.42%	4475.63	428	24961	3211.00	9528.90	10250.90	13385.36	75.83	57414.85	16.08
Sync	4500	19	0.42%	4475.63	428	24961	3211.00	9528.90	10250.90	13385.36	75.83	57414.85	16.08

Requests	Executions			Response Times (ms)							Throughput	Network (KB/sec)	
Label	#Samples	FAIL	Error %	Average	Min	Max	Median	90th pct	95th pct	99th pct	Transactions/s	Received	Sent
Total	4500	49	1.09%	7526.70	134	249254	6166.50	17408.20	22149.65	25108.06	15.91	11966.40	3.16
Async	4500	49	1.09%	7526.70	134	249254	6166.50	17408.20	22149.65	25108.06	15.91	11966.40	3.16

Normal Threads: Thread, and ThreadPool

Thread = OS Thread. Mapping 1:1 to OS Threads

```
Thread t01 = new Thread(() => Console.WriteLine($"Hello  
{Thread.CurrentThread.ManagedThreadId}"));  
t01.Start();
```

ThreadPool = pool of threads

```
ThreadPool.QueueUserWorkItem
```

Important .Abort, .Kill –deprecated!

Use CancellationTokenSource and CancellationToken
(or kill new PROCESS)



Demo – 01 Thread and ThreadPool and...

Demo01AThreads

Demo01BThreadPool

Threads/Task/Pool - [app].runtimeconfig.json | Env

System.GC.Server: (WS – one CPU, true, SRV – dedicated threads)

System.GC.Concurrent: (background), one per logical vCPU

System.GC.HeapAffinitizeMask, System.GC.CpuGroup, System.GC.HeapCount,
System.GC.HeapHardLimit, LOH | SOH | POH Settings, System.GC.Name (custom),
System.GC.ConserveMemory

DOTNET_Thread_UseAllCpuGroups | DOTNET_Thread_AssignCpuGroups

System.Threading.ThreadPool.MinThreads | System.Threading.ThreadPool.MaxThreads

System.Threading.ThreadPool.UseWindowsThreadPool

System.Threading.ThreadPool.Blocking.* : sync-over-async case & block, AVOID (in general)

Also: AppContext.SetSwitch (some settings)

Ps. COMPlus_ => DOTNET_

Task Parallelism, Data Parallelism: Task Parallel Library

Why “threads” are BAD?

Does not scale as well as you think. Thread/Context switching over head

Thread starvation. You own synchronization

TPL handles: partitioning, the scheduling of threads on the ThreadPool, cancellation support, state management, ...

Also: Concurrent collections – to make life a little bit simpler...

Parallel.For / ForEach

PLINQ: 2 syntaxes – similar to SQL and adopting fluent API approach

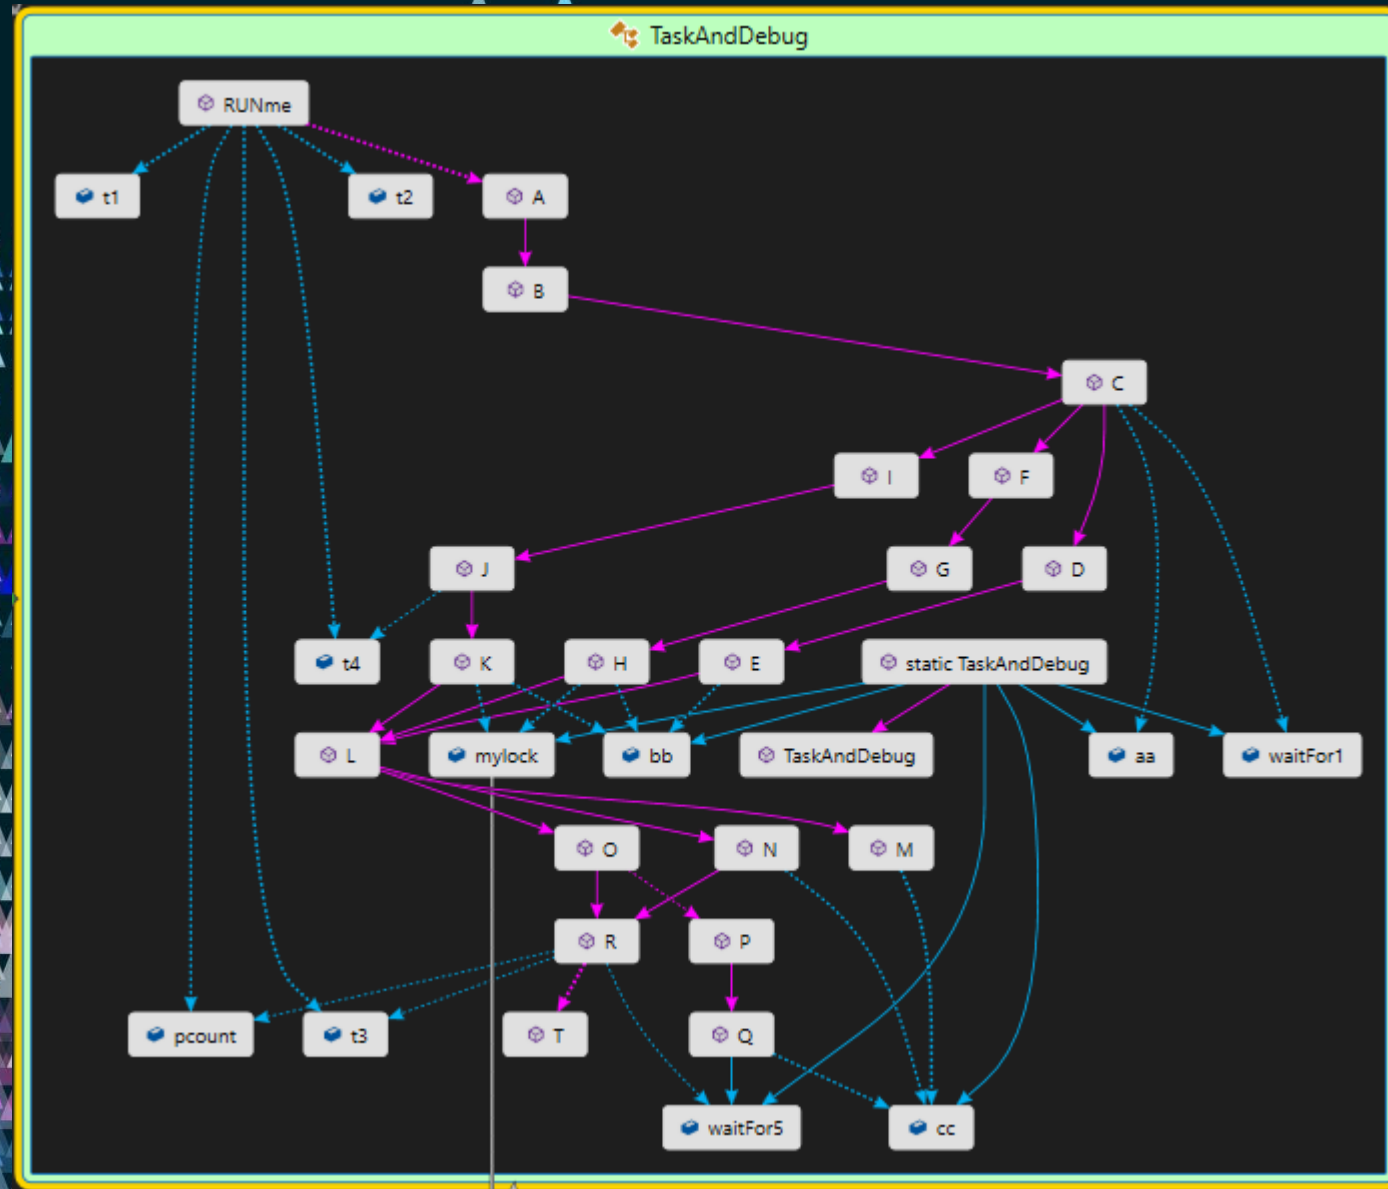
async / await - for – I/O Bound work, like:

HttpClient, Disk I/O, Socket I/O. Also: Task.Run !

Not (always) new thread. Thread Context – the same (or different) :

`taskA.ConfigureAwait(true)`

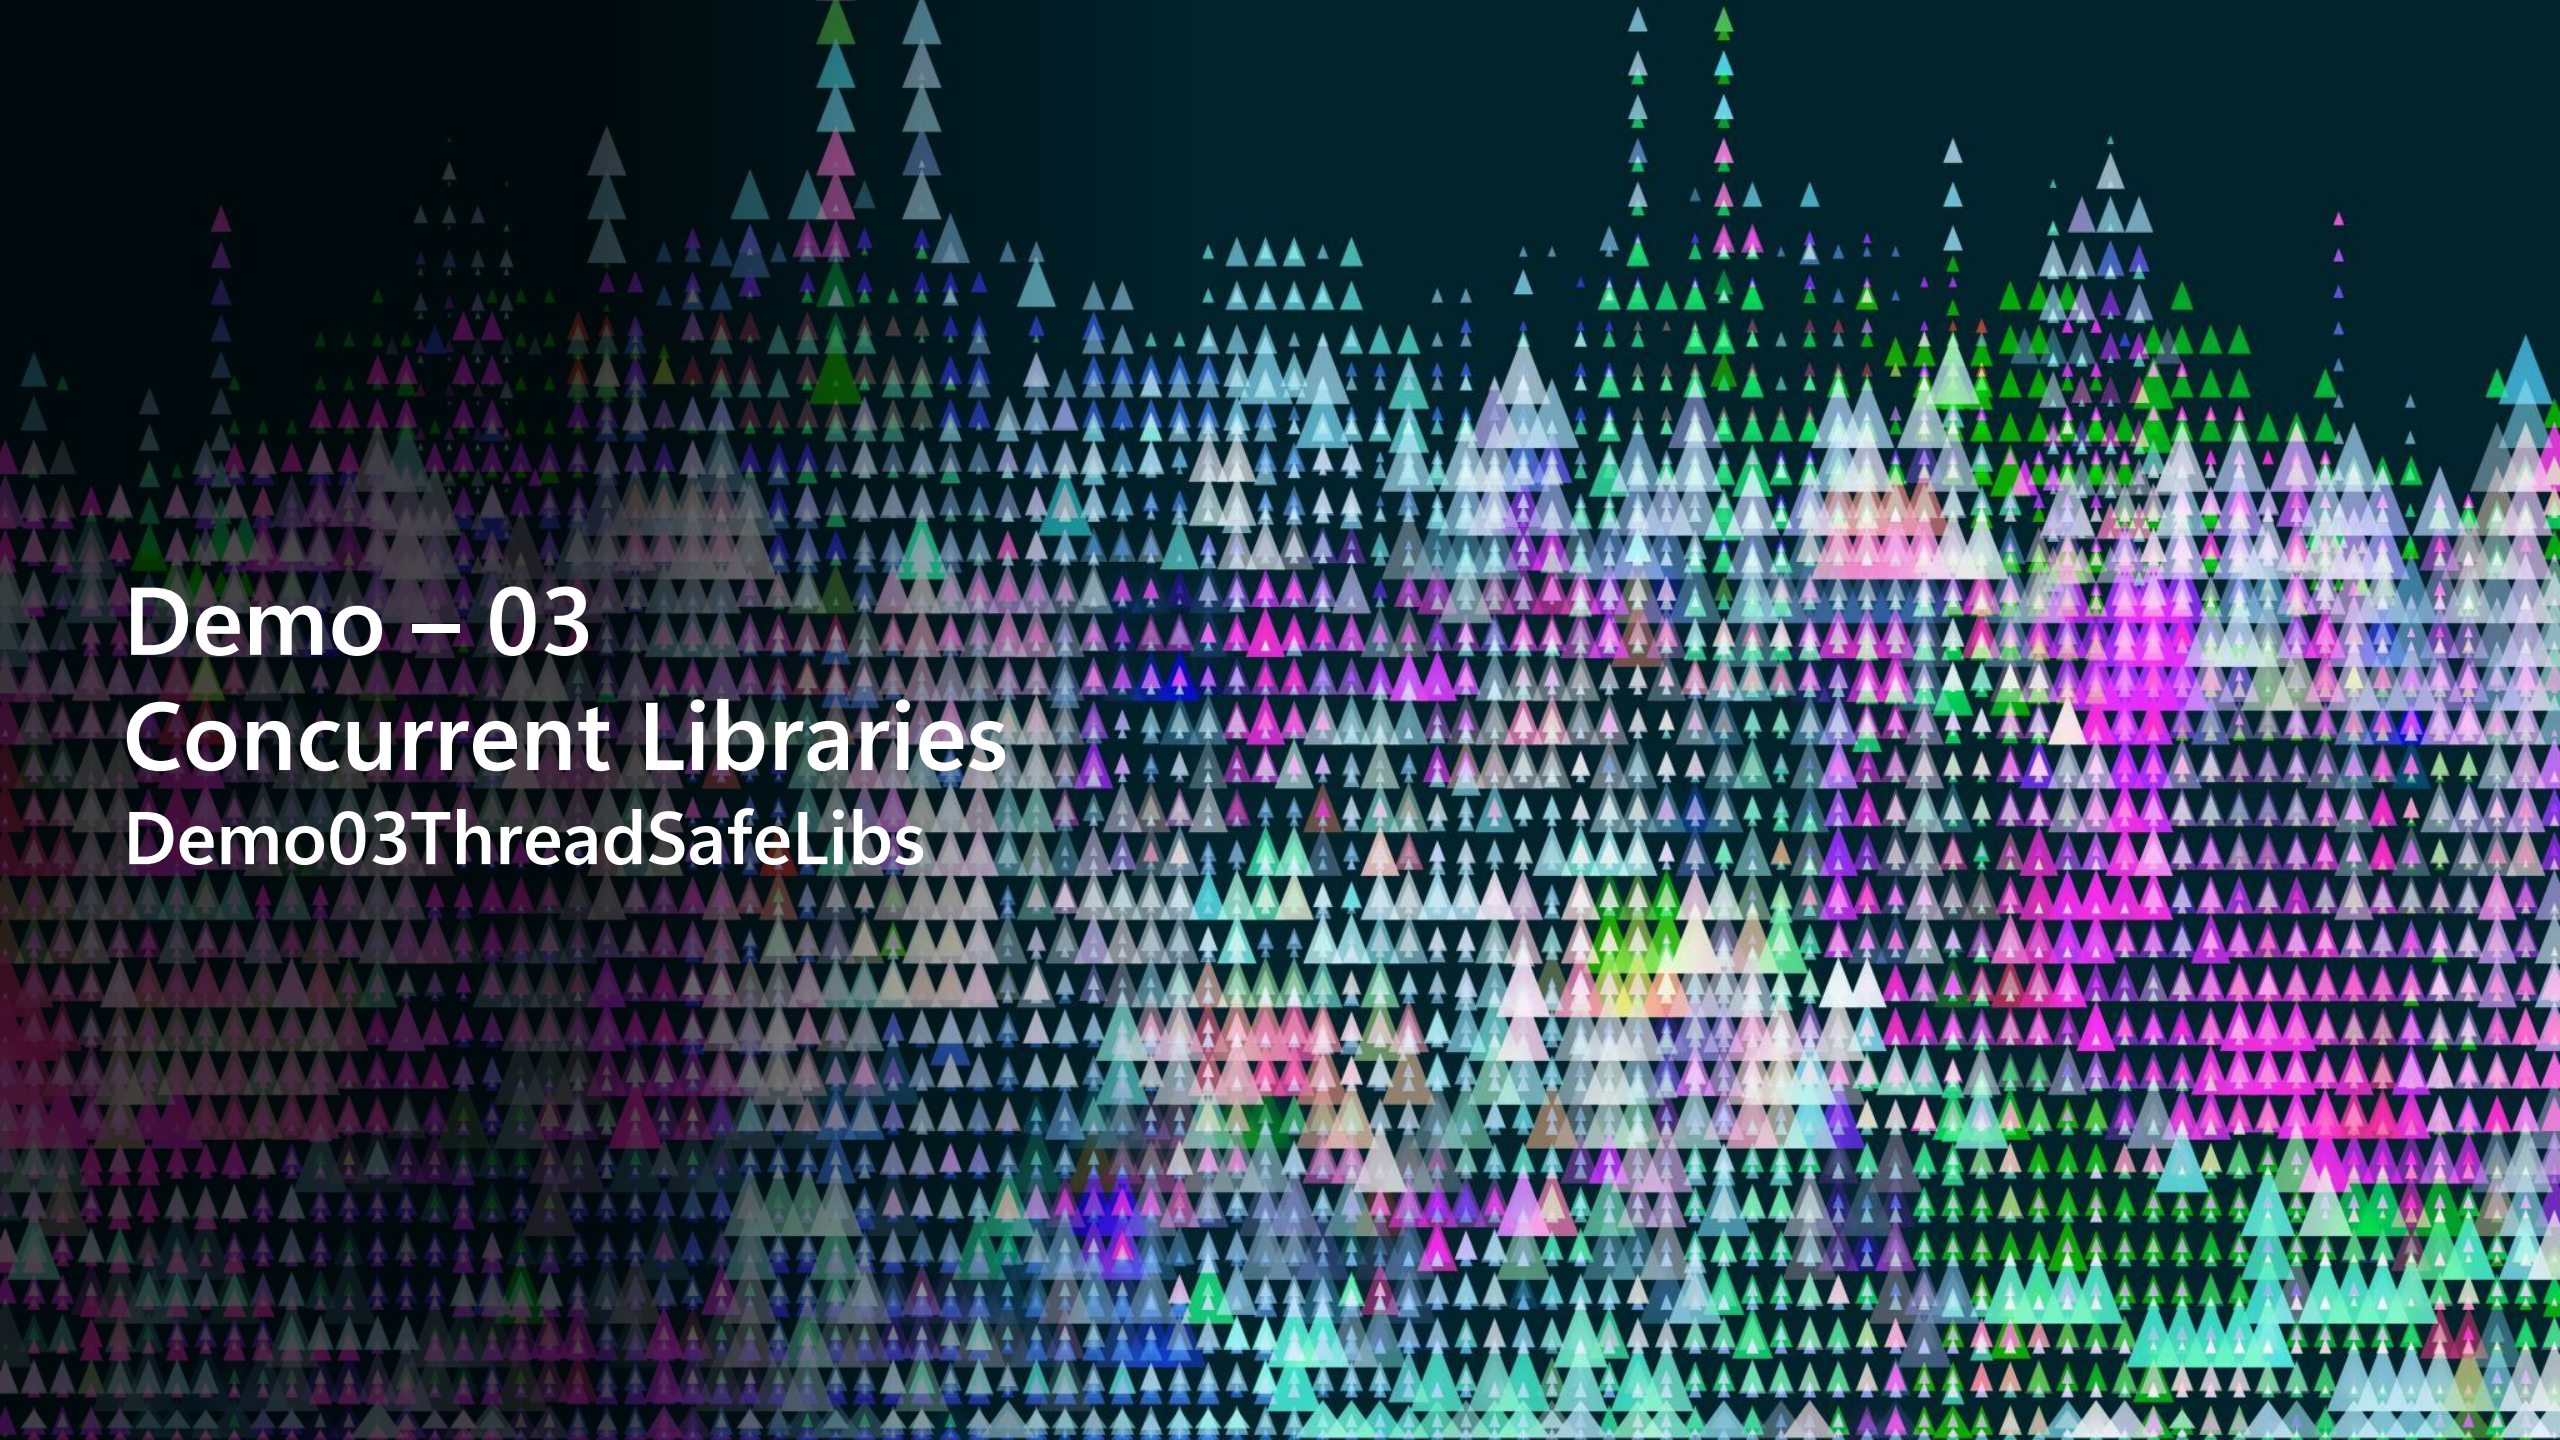
Demo02TPLSynchronization





Case – Blazor

Server and client-side rendering
Blazor-SimplestDemo

The background of the slide is a dense, colorful pattern of small triangles. The triangles are arranged in a way that creates a sense of depth and movement, with colors ranging from dark blues and purples to bright greens, yellows, and pinks. The pattern is more concentrated on the right side and fades out towards the left.

Demo – 03

Concurrent Libraries

Demo03ThreadSafeLibs



Demo – 04 PLINQ

Demo04PLINQ

Demo04PLINQSetPixel

TPL Dataflow

Message passing and parallelizing CPU-intensive **AND** I/O intensive application.
High throughput **AND** low latency.

Concept like actor-based programming (Orleans, Akka)

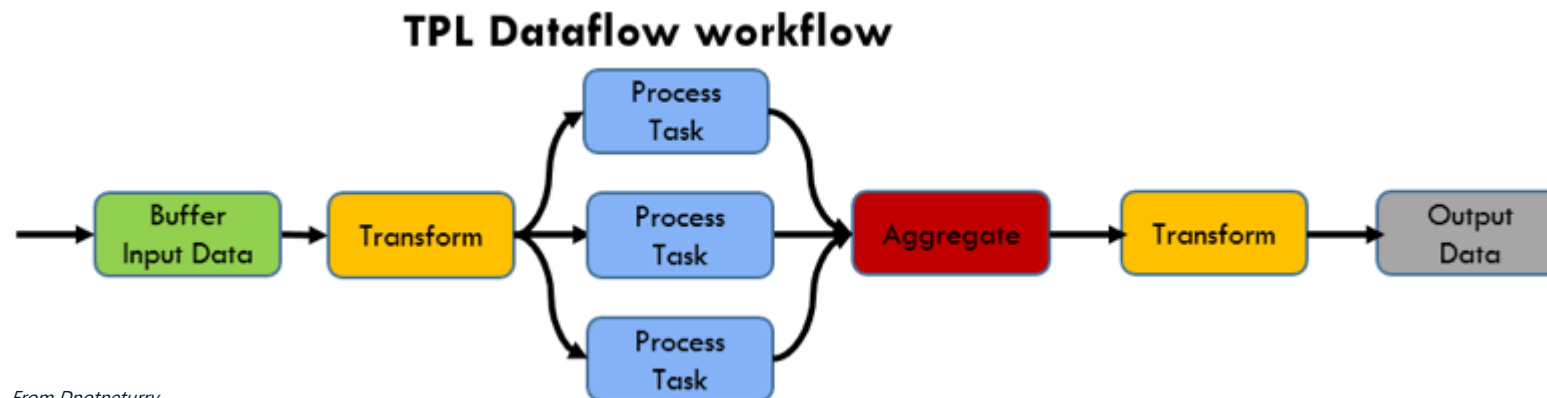
In-process message passing, ETL like pipelines, not distributed

VERY similar semantic

Helpful to avoid deadlocks

Ps. Dataflows are also in Power Query – very similar semantic

Blocks: Source, Propagator, Target, Pipeline





Demo – 05 Data Flow

Demo05DataFlow

Rx – Reactive Extension for .NET

.NET library for processing event streams – [here](#). Observable pattern

Better Events, modern Streams...

```
IObservable bigTrades = trades.Where(trade => trade.Volume > 1_000_000);  
bigTrades.Subscribe( ...)
```

(Like IEnumerable, but will add a new element when available)

Fully integrated with ecosystem: combination with TPL Dataflow, (P)LINQ etc

Scenarios:

Complex Event Processing, **CQRS**, Integration with IoT Brokers, UI!, Tracking in **real time** (Web, GPS, ...). Domain Model Events. Broadcasting.

Event passing vs complex thread topologies

Also, [REAQTOR](#) (Rx + state, durability)

Also, integration with Blazor (and others UI frameworks!)

Also (project from 2010) many additional integrations...

The background of the slide is a dense, abstract pattern of small triangles. The triangles are arranged in horizontal bands of varying colors, including shades of blue, green, purple, pink, and yellow. The overall effect is a vibrant, pixelated or mosaic-like texture that fills the entire frame.

Demo – 06 Reactive Extensions

Demo06RxNet

Summary – how to “think”

Use mix of Data and Task Parallelism (map-reduce, fork-join etc)

Mandatory Data Parallelism for larger sets

Use Task Parallel Library – “serialize” parallel code

async/await for all I/O bound. Actors “like” approach – maybe DataFlow | Reactive Extensions

Think before “rewriting”

Treat “TPL” as a **standard** toolbox

Classic threads – only for exams/tests/jobs interviews ;) or – libraries. Time to **write tests!**

Prepare code – avoid “globals”, keep local calculation etc...

Supportive materials:

.NET docs - <https://github.com/dotnet/docs>

Excellent - <https://github.com/dotnet/runtime/tree/main/docs>

<https://github.com/dotnet/performance>

■ FEEDBACK

.NET - Programowanie równoległe i trochę rozproszone - kiedy (i jak)



Tomasz Kopacz

<https://mstechsummit.pl/user.html#!/lecture/MSTS24-19c1/rate>

Questions?

tkopacz@microsoft.com

Demos: <https://github.com/tkopacz/msts2024-net-parallel>

